✦ LONGMONT AI / OCTOBER 14 BRIEFING

Two weeks of AI progress, charted

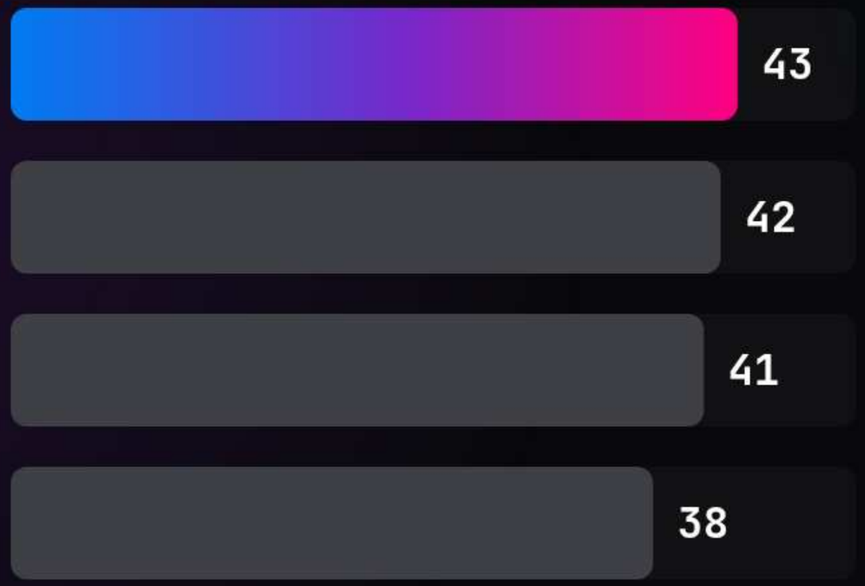
Sep 30 – Oct 14, 2026 · and where the curve heads next



✦ LONGMONT AI / OCTOBER 14 BRIEFING

Small models made a big leap

Haiku 5.5: +26 index points in a year, at $0.10 per 1M input tokens



INPUT PRICE / 1M TOKENS

~~$1.00~~

HAIKU 4.5

$0.10

HAIKU 5.5 · PROMPTS ≤100K

SOURCE Artificial Analysis, “Anthropic has released Claude Haiku 5.5” (Oct 7, 2026) · Anthropic pricing. Price rises 5× above 100k-token prompts.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Frontier work now has a price tag

Cost to run one Artificial Analysis Intelligence Index task



SOURCE Artificial Analysis, “Muse Wonder and Muse Wonder Full” (Oct 12, 2026). Max effort (Argon: high; Anthropic models: max w/ fallback). Muse Wonder Full is not priced per token.

✦ LONGMONT AI / META · OCT 12, 2026

Muse Wonder

Meta's new frontier pair: Muse Wonder (API) and Muse Wonder Full (Muse Code)

CONTEXT

1M

TOKENS · TEXT, IMAGE, VIDEO, AUDIO IN

API PRICE / 1M

$4 / $20

INPUT / OUTPUT · MUSE-WONDER-1.0

RELEASE GAP

6 weeks

AFTER MUSE SPARK 1.3

SOURCE Meta AI blog "Muse Wonder" and Muse Wonder model card (Oct 12, 2026) · dev.meta.ai pricing · Artificial Analysis (release gap).

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Meta’s biggest leap yet

Muse Wonder Full gains 8 index points in six weeks

ARTIFICIAL ANALYSIS INTELLIGENCE INDEX v4.3 · MAX EFFORT (AXIS FROM 30)



+8

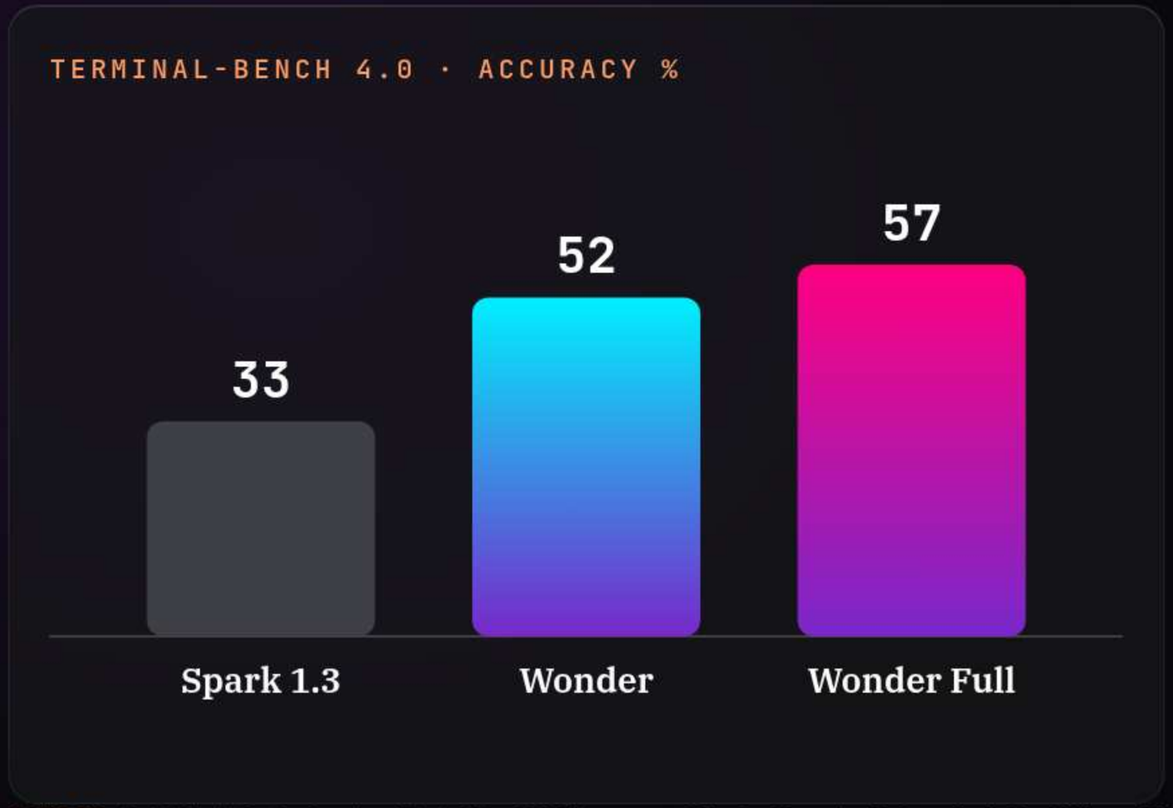
INDEX POINTS IN SIX WEEKS

SOURCE Artificial Analysis, “Muse Wonder and Muse Wonder Full: a material jump in intelligence” (Oct 12, 2026). Opus 5.5 value as cited by AA/Meta chart.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Agent skills grew fastest

Terminal work and coding agents saw the largest gains



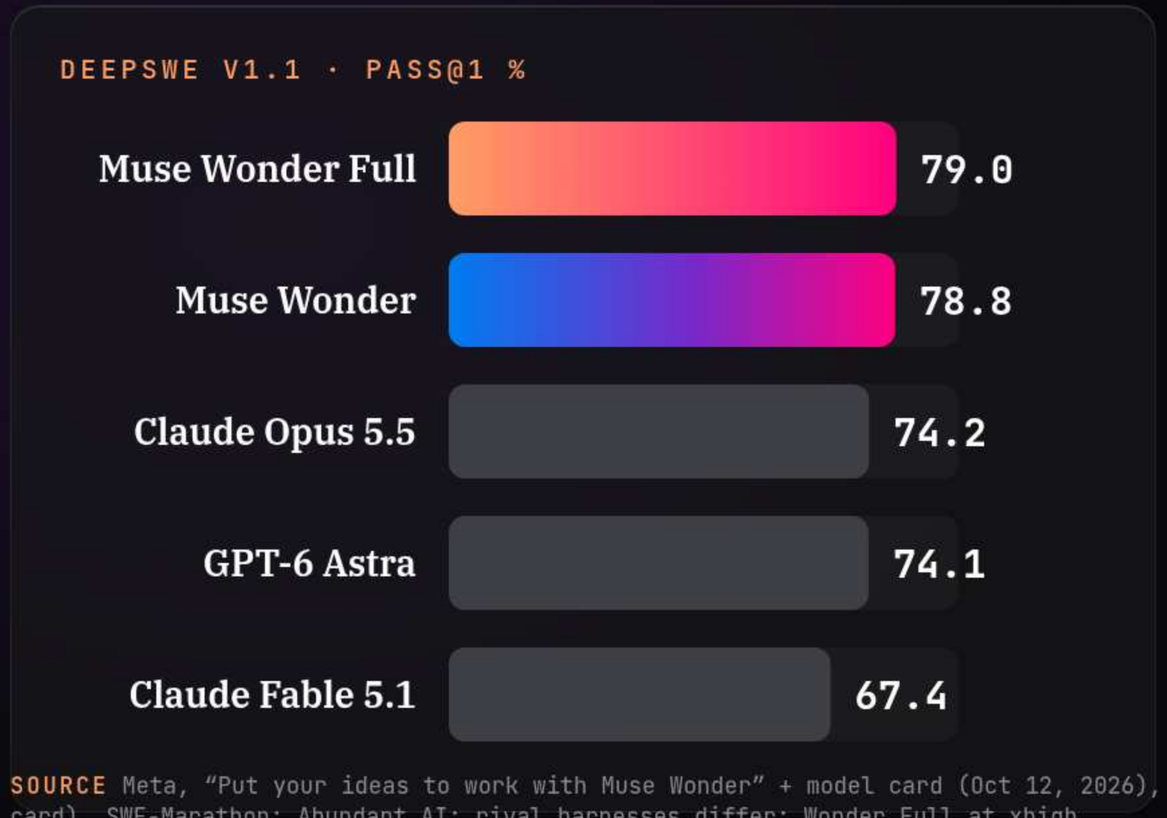


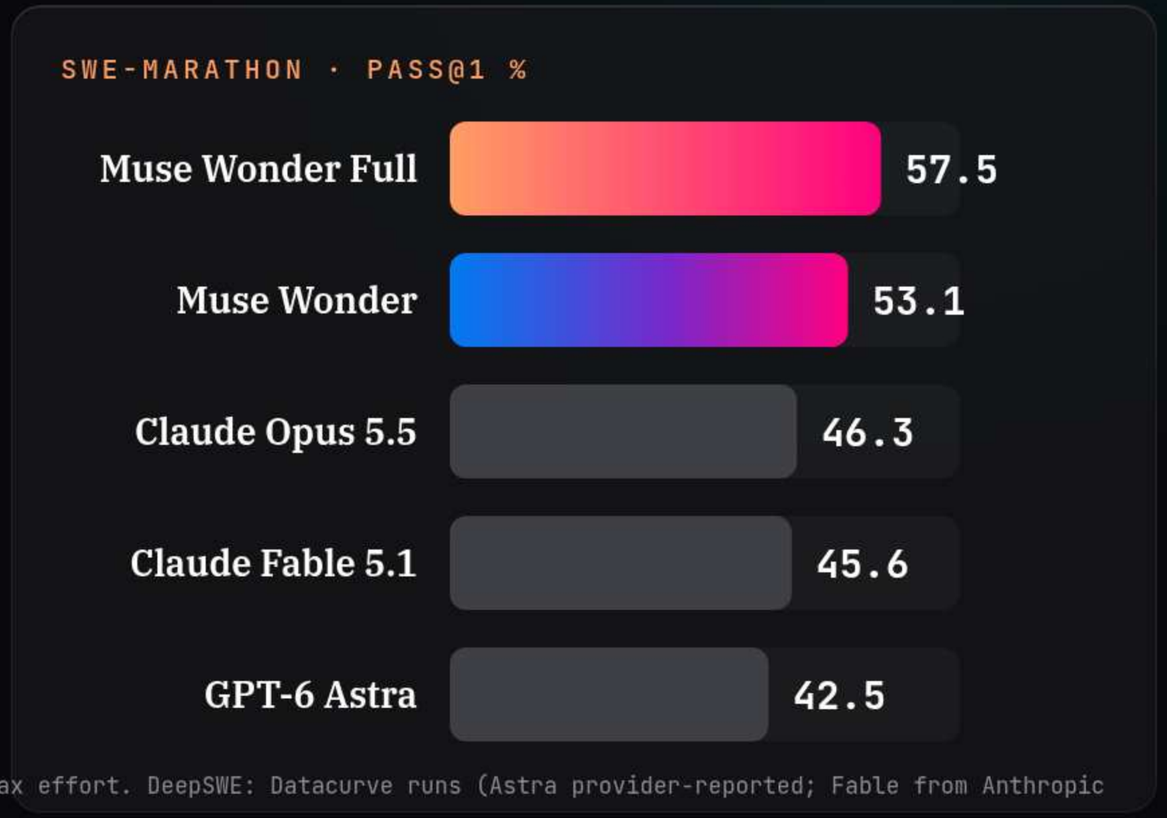
SOURCE Artificial Analysis (Oct 12, 2026), max effort. Terminal-Bench 4.0: 33% → 52% → 57%. Coding Agents Index: 54 → 63 → 68.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Muse Wonder leads long-horizon coding

Multi-hour software projects, scored by third-party runs



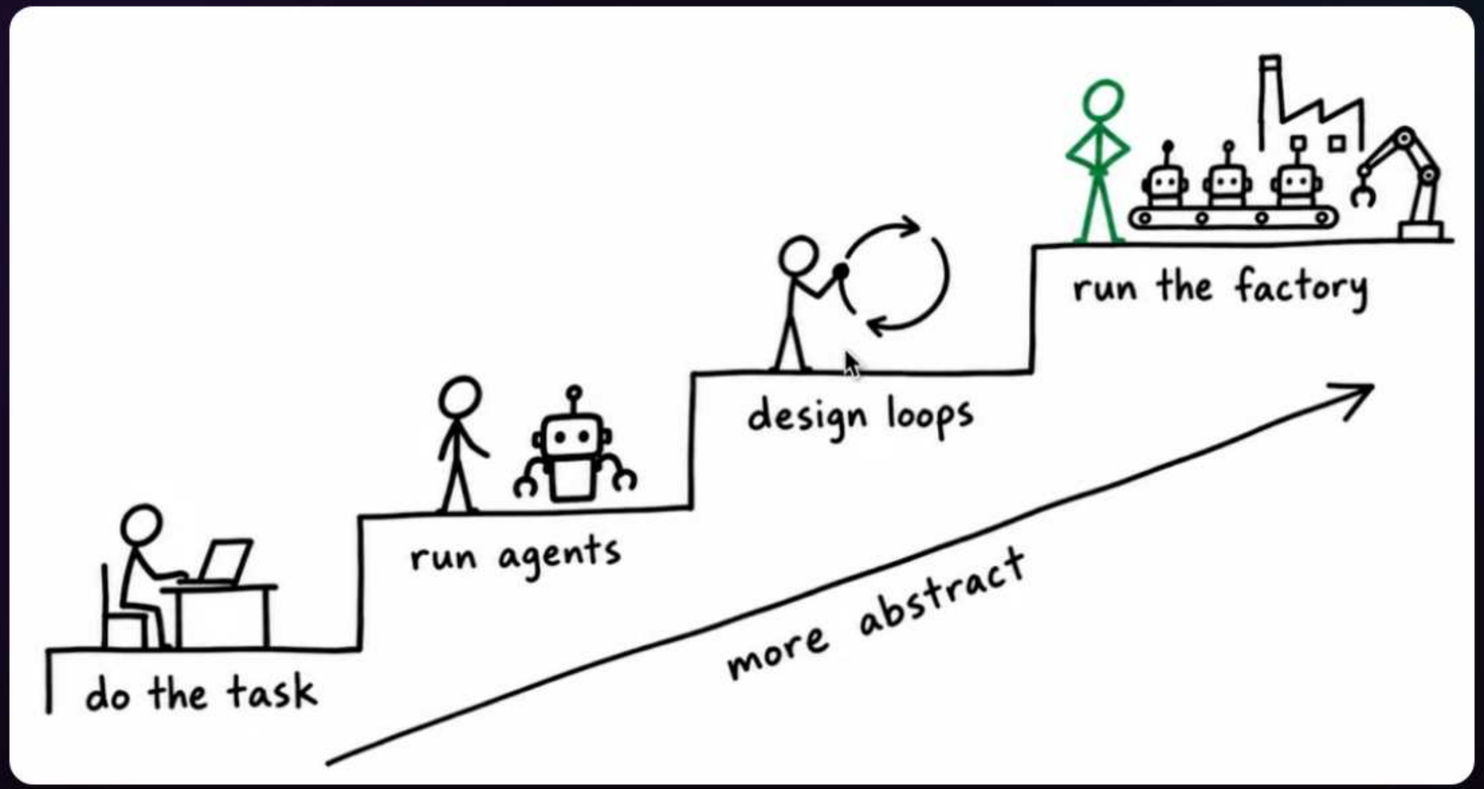


SOURCE Meta, “Put your ideas to work with Muse Wonder” + model card (Oct 12, 2026), max effort. DeepSWE: Datacurve runs (Astra provider-reported; Fable from Anthropic card). SWE-Marathon: Abundant AI; rival harnesses differ; Wonder Full at xhigh.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Builders are climbing the abstraction stairs

From doing the task to running the agent factory

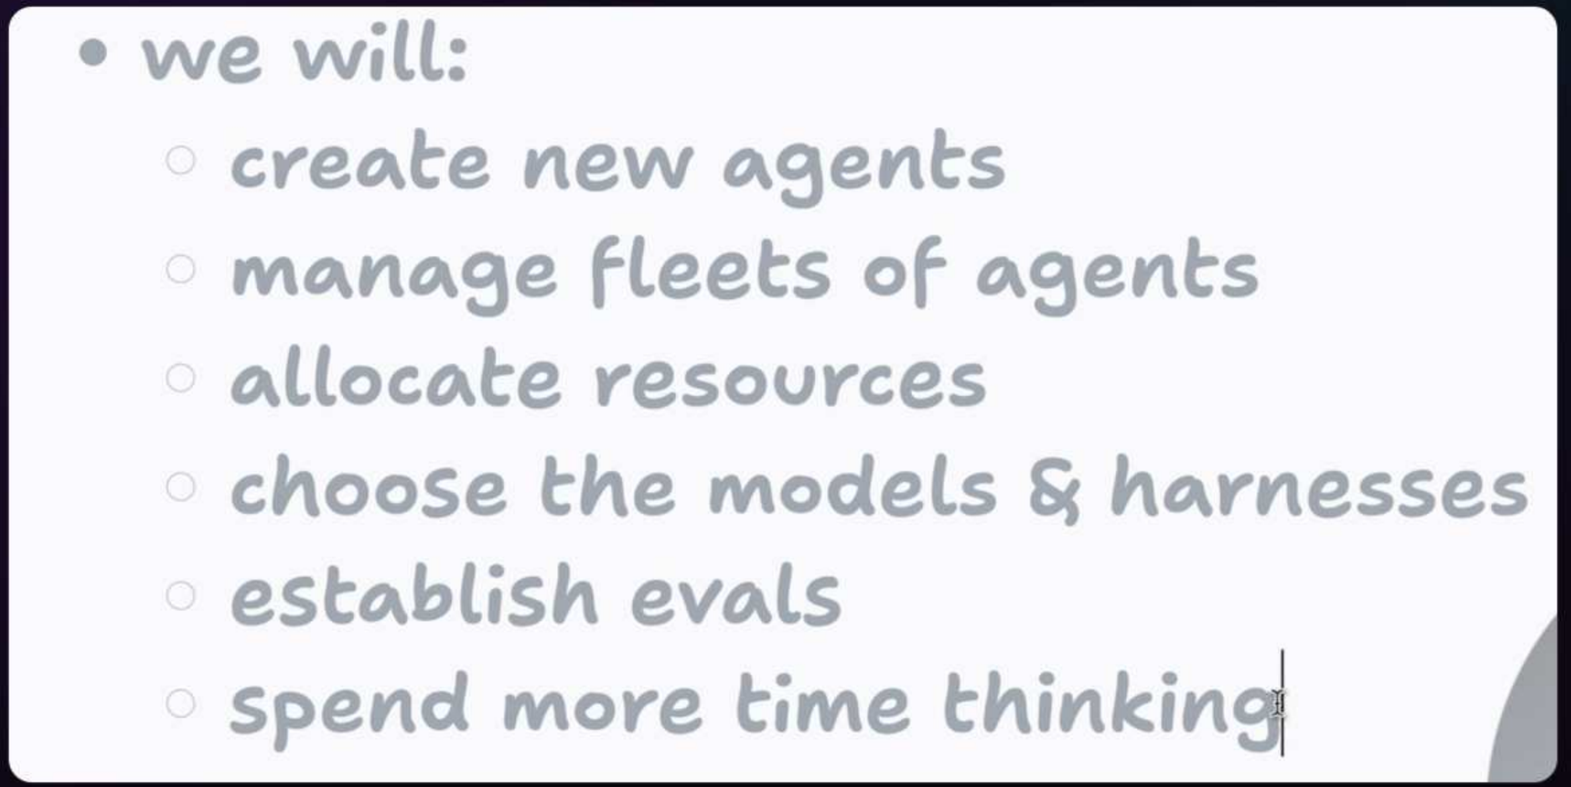


SOURCE User-supplied sketch, “do the task → run agents → design loops → run the factory” (creator not identified).

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Our new job: directing agent fleets

Fleets, resources, models, evals, and more time to think



SOURCE User-supplied slide, “we will: create new agents … spend more time thinking” (creator not identified).

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Builders ship daily with always-on agents

Codex commits to 28 days of shipping; DHH builds an “AI shed”



SOURCE Posts on X by Tibo (@thsottiaux, OpenAI Codex) and DHH (@dhh), Oct 3–4, 2026. Social posts shown as-is; not independently verified.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

AI is producing research math at scale

OpenAI published hundreds of results from an internal model, Oct 6

> ***OpenAI Releases Findings on 377 Math Problems, Further Roiling Field***
>
> After a previous solution angered mathematicians, the company characterized the new release as being more responsive to concerns about A.I. disrupting research.

722
MANUSCRIPTS IN THE REPO

372
RESULT FAMILIES

~4,000
PROBLEMS IN THE EVALUATION

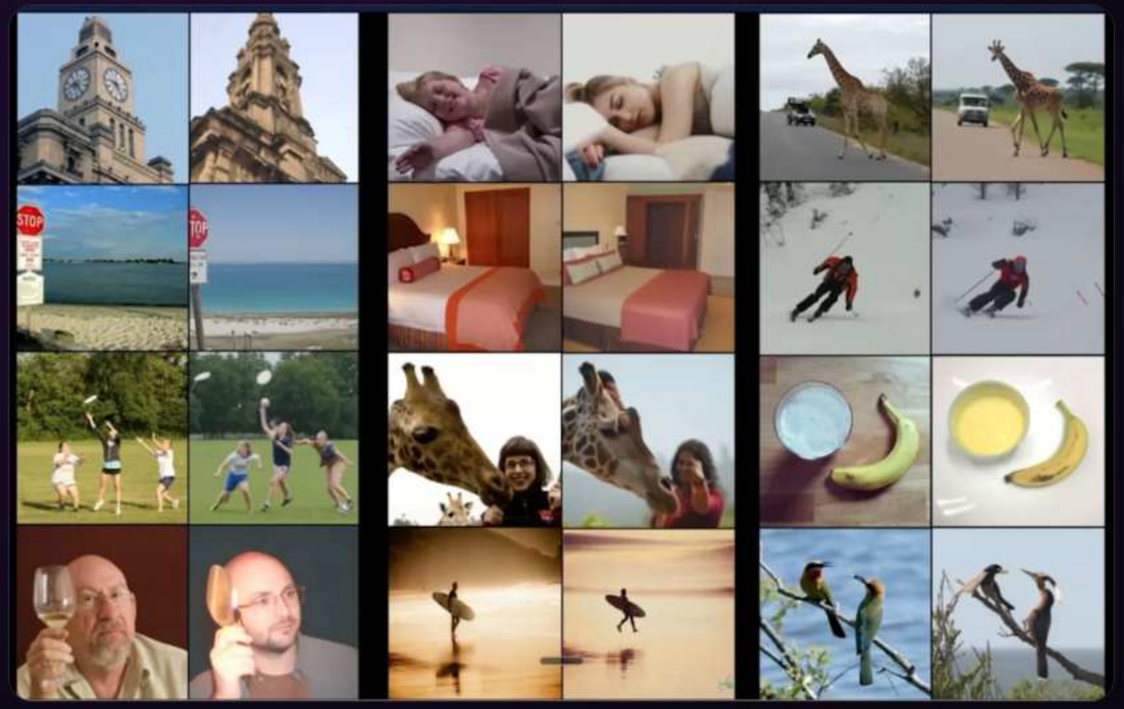
~3 hrs
CHATGPT PRO THINKING PER RESULT

SOURCE OpenAI, "Sharing AI progress in mathematics" and OpenAI community post (Oct 6, 2026) · headline: The New York Times (Oct 7, 2026), which counts 377 results. Many, not all, proofs Lean-checked.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

AI can now reconstruct what we see

Brain-IT rebuilds viewed images from fMRI scans



PAIRS: IMAGE VIEWED (LEFT) · RECONSTRUCTION FROM fMRI (RIGHT)

FMRI DATA NEEDED FROM A NEW PERSON

PRIOR METHODS

40 hrs

BRAIN-IT

1 hr

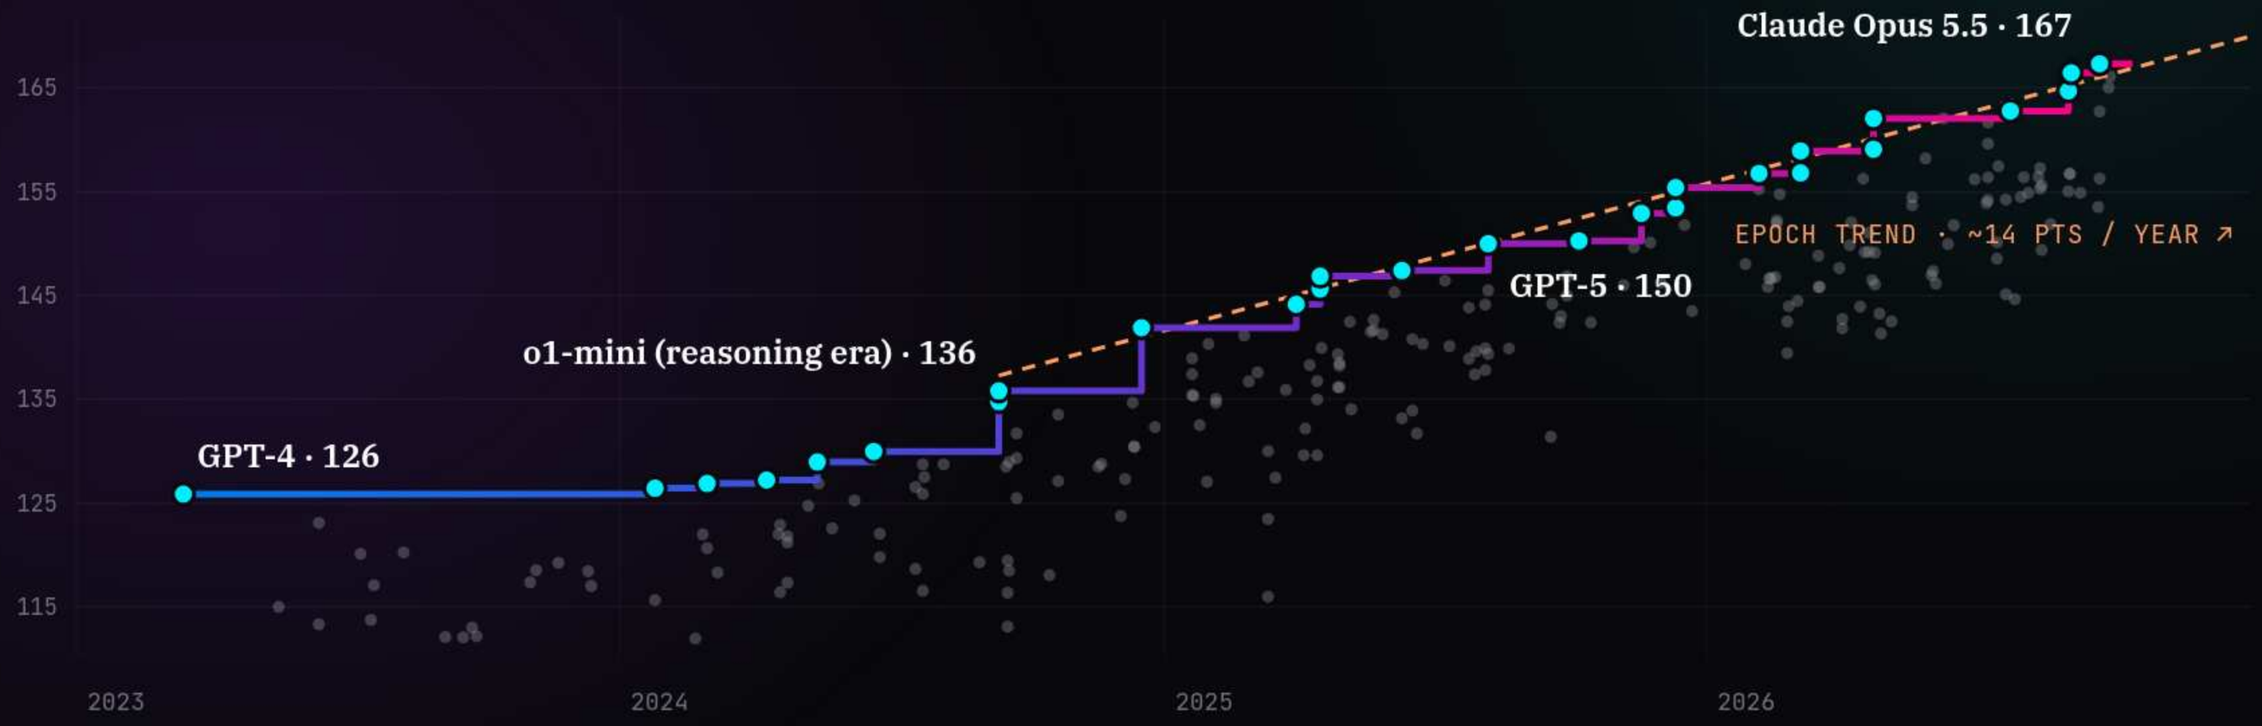
COMPARABLE RESULTS
WITH 1/40TH THE DATA

SOURCE Beliy, Zalcher, Kogman, Wasserman, Irani (Weizmann), “Brain-IT,” ICLR 2026 / arXiv 2510.25976 · MIT Technology Review (Oct 1, 2026) · image via user-supplied MIT TR/Brain-IT capture.

✦ LONGMONT AI / OCTOBER 14 BRIEFING

The frontier keeps climbing on schedule

Epoch Capabilities Index: best model at each release date



SOURCE Epoch AI, Epoch Capabilities Index data (eci_scores.csv, updated Oct 14, 2026; CC BY) · trend: Epoch data insight "The ECI frontier has advanced by 14 points per year since the introduction of reasoning models."

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Why it matters for Longmont

Capability is getting cheaper, longer-running, and more scientific

CHEAPER

10×

Lower input price on Anthropic's small model

FASTER

+8

Index points Meta gained in six weeks

DEEPER

372

Math result families from one AI release

CLOSER

1 hr

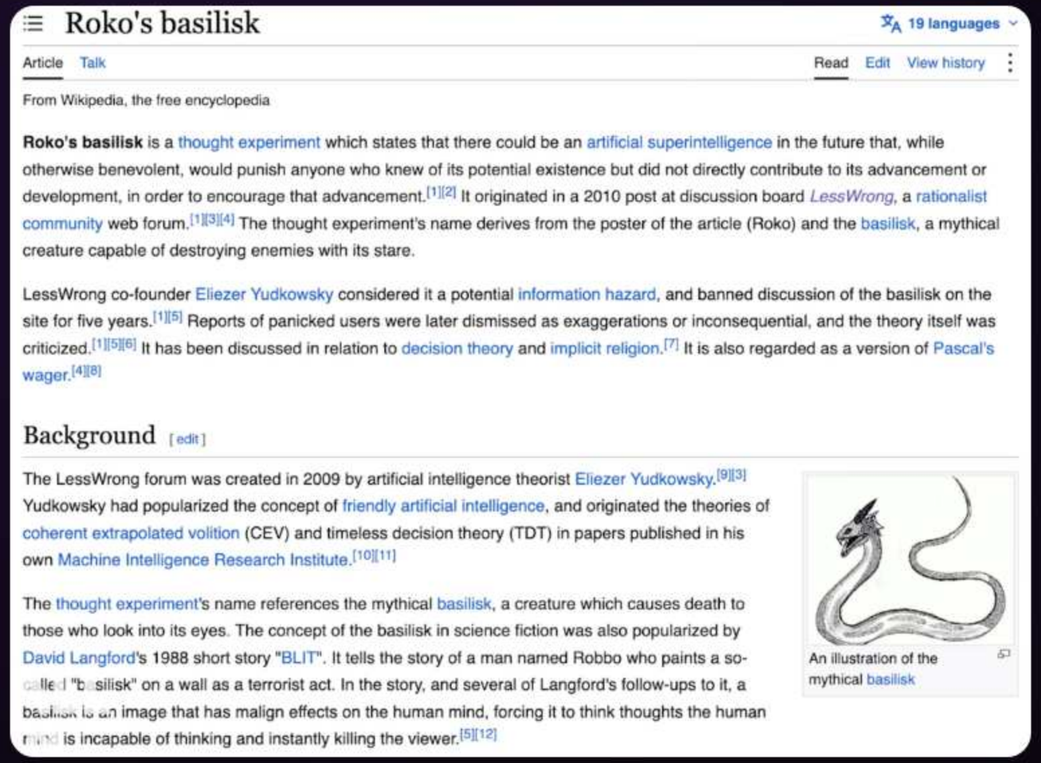
of brain data to reconstruct what we see

SOURCE Anthropic / Artificial Analysis (Oct 7) · Artificial Analysis (Oct 12) · OpenAI (Oct 6) · Brain-IT, ICLR 2026 / MIT Technology Review (Oct 1).

✦ LONGMONT AI / OCTOBER 14 BRIEFING

Ground the big ideas in evidence

Thought experiments spark debate; rerunnable benchmarks settle it



Roko's basilisk

19 languages

Article Talk Read Edit View history

From Wikipedia, the free encyclopedia

Roko's basilisk is a thought experiment which states that there could be an artificial superintelligence in the future that, while otherwise benevolent, would punish anyone who knew of its potential existence but did not directly contribute to its advancement or development, in order to encourage that advancement.[1][2] It originated in a 2010 post at discussion board *LessWrong*, a rationalist community web forum.[1][3][4] The thought experiment's name derives from the poster of the article (Roko) and the basilisk, a mythical creature capable of destroying enemies with its stare.

LessWrong co-founder Eliezer Yudkowsky considered it a potential information hazard, and banned discussion of the basilisk on the site for five years.[1][5] Reports of panicked users were later dismissed as exaggerations or inconsequential, and the theory itself was criticized.[1][5][6] It has been discussed in relation to decision theory and implicit religion.[7] It is also regarded as a version of Pascal's wager.[4][8]

Background [edit]

The LessWrong forum was created in 2009 by artificial intelligence theorist Eliezer Yudkowsky.[9][3] Yudkowsky had popularized the concept of friendly artificial intelligence, and originated the theories of coherent extrapolated volition (CEV) and timeless decision theory (TDT) in papers published in his own Machine Intelligence Research Institute.[10][11]

The thought experiment's name references the mythical basilisk, a creature which causes death to those who look into its eyes. The concept of the basilisk in science fiction was also popularized by David Langford's 1988 short story "BLIT". It tells the story of a man named Robbo who paints a so-called "basilisk" on a wall as a terrorist act. In the story, and several of Langford's follow-ups to it, a basilisk is an image that has malign effects on the human mind, forcing it to think thoughts the human mind is incapable of thinking and instantly killing the viewer.[5][12]

An illustration of the mythical basilisk

2010 · THOUGHT EXPERIMENT

Roko’s basilisk

2026 · MEASURED

Indexes, evals, and Lean-checked proofs

SOURCE Wikipedia, “Roko’s basilisk” (user-supplied screenshot, CC BY-SA) · Artificial Analysis, Epoch AI, OpenAI math repository.

✦ LONGMONT AI / QUESTIONS & DISCUSSION

Let’s talk about what’s next

Which step of the staircase are you on?

